

Advancing Layer Surface Mesh Generation

Jasmeet Singh* and Carl Ollivier-Gooch†

University of British Columbia, Vancouver, British Columbia, V6T1Z4

A new method of generating advancing layer surface meshes is demonstrated. We generate semi structured quad-dominant meshes with the ability to have local control over the aspect ratio. Semi-structured two-dimensional anisotropic meshes in the boundary layer regions have been shown to have superior fluid flow simulation results. However, the discretization of the surfaces poses challenges in replicating the same for volume meshes. Our surface mesh is constructed by independently meshing surface segments using an advancing layer methodology. Point placement in layers and local reconnection helps produce a valid surface mesh at each step of mesh generation. We demonstrate the ability of the meshing algorithm to tackle fairly complex geometries and coarse initial surface discretization.

I. Introduction

THE need to generate advancing layer meshes is prevalent in mechanical engineering and aerospace applications [1]. Boundary layer meshes generated by advancing layer strategy have found applications even in biomedical computing [2]. Wherever there is a need to refine the elements around the boundary of an object, a high-aspect ratio advancing layer mesh is desired for carrying out the simulations. Such a mesh helps in capturing the gradients normal to the direction of a specified boundary without refining the region near the boundary. Hence, computational cost and memory usage are considerably reduced.

In many aerodynamic applications, a boundary layer volume mesh with anisotropy in specific directions is very much desired. Some of the initial attempts at generating stretched element meshes in two-dimensions used a Delaunay mesh and a locally mapped space to get the required anisotropy [3, 4]. Other techniques either used an approach of using a locally structured or semi structured mesh for the regions requiring high anisotropy[5, 6] or using an advancing layer strategy to generate the mesh in layers while terminating locally whenever necessary([7–9]). Several techniques have also been developed to generate boundary layer volume meshes [10–14].

The methods stated above work on simplicial meshes, ie meshes with triangles in two dimensions and tetrahedra in three dimensions. These simplicial meshes have an advantage of good flexibility in discretizing the domain, especially if it contains complex features. However, the vertex connectivity of simplicial mesh elements is high. The number of edges in a tetrahedral mesh is seven times the number of vertices. On the other hand, in a hexahedral mesh, the number of edges is only three times the number of vertices(asymptotically). The cost of vertex-based discretization is directly proportional to the number of edges in a mesh. Hence, non-simplicial meshes are substantially more efficient than simplicial meshes for a given number of unknowns or grid points. Also, regular arrays of nonsimplicial elements may also enhance accuracy, owing to a local cancellation of truncation errors that may not occur on groups of nonsimilar simplicial elements[15]. Quadrilateral elements have been preferred over triangles in highly stretched two-dimensional grids due to their lower connectivity [16]. The advantages of non-simplicial mesh elements have resulted in fully non-simplicial mesh generation techniques [17, 18]. However, simplicial mesh elements are necessary to deal with the complexity in the geometry. Hence, a hybrid-mesh containing both simplicial and non-simplicial mesh elements is a reasonable choice for mesh generation. We generate a surface mesh which is quad dominant with some triangles in it. The present draft shows meshes in which quads are represented by a pair of almost right angled triangles. The combination of these pair of triangles to quad elements is a relatively straightforward process and would be included in the final paper.

Stretched volume meshes are generated usually from an initial triangulation of a surface. These surfaces are usually represented by NURBS in various CAD packages. Generation of the surface mesh requires a separate methodology or algorithm as compared to the two-dimensional mesh generation as the mesh elements include a third dimension. The surface mesh can be created through a parametric mapping from two-dimensional mesh or by a direct three-dimensional method. In the first method, the mesh can be created by mapping a two-dimensional delaunay mesh on a surface[19, 20] using surface parametrization. This method, however, doesn't always form well shaped surface elements if the surface

*Graduate Research Assistant, Mechanical Engineering, jasmeet.singh.mec11@iitbhu.ac.in

†Professor, Mechanical Engineering, cogg@mech.ubc.ca

derivatives vary widely over the domain [21]. Also, these methods do not intend to generate an anisotropic surface mesh. Equivalent method has also been introduced to generate the surface mesh using advancing front technique by utilizing a metric derived from the first fundamental form of the surface [22, 23].

Another way of creating a surface mesh is by creating mesh elements directly on the surface geometry. Lao and Lu[24] have formulated an advancing front approach for arbitrary three dimensional surfaces. This approach creates triangular elements on the surface geometry. It uses surface normal, tangent and projection computations to generate the surface mesh. In another work [25], they use this triangular mesh to generate an all quad mesh. For this method to work successfully, the underlying triangular mesh needs to be sufficiently well resolved in geometrically complex features of the mesh. Also, meshes developed by this method also do not involve anisotropic surface mesh elements.

We present a surface mesh algorithm that creates a surface discretization with required directional anisotropy and mostly quad elements. Such a surface mesh can be a stepping stone towards an anisotropic hex-dominant boundary layer volume mesh. The surface mesh generation algorithm is based on a closed advancing front method. The method needs minimal user input and automatically generates an advancing layer mesh from a given input surface triangulation. The input surface is first segmented into several sub-surfaces so that each of them can be meshed independently. Boundaries of the surfaces are identified and the points on the boundary iteratively march towards the interior of the surface. Initial extrusion length and growth ratio can be defined to give the required level of anisotropy at each point of the mesh or for a given boundary. A valid surface mesh is produced after each layer which gives the user freedom to advance until a level of anisotropy or until the advancing front routine terminates.

However important may the advancing layer mesh be, its creation poses challenges for geometries that are complicated. Sharp corners and highly curved regions in the input geometry pose robustness issues for most mesh generators. Our advancing front routine includes several validation checks to avoid these issues, however, we are working on making it more robust for complex geometries.

II. Methodology

The surface mesh algorithm presented in this paper starts by importing a surface triangulation of and creating an interpolated model of the surface. Steps taken from importing surface triangulation to the mesh completion are detailed in the following subsections.

A. Surface Import and Segmentation

The surface mesh generator reads a surface triangulation as a faceted geometry. It then uses the Common Geometry Module (CGM)[26] facet based geometry kernel to create an interpolated geometry from the facet based triangulation. G^1 continuous curves and surfaces are created from triangles which form quartic Bezier triangle patches. The geometry imported by CGM is used as the background surface to create the mesh. Robust and cost efficient geometry and topology query operations in CGM helps us in creating the surface mesh. The most fundamental requirements of the surface mesh from the geometry are - closest point or point projection routine and surface normal evaluations. These operations are frequently used to project point onto the underlying geometry, compare normal of a triangle to the normal at a point on the underlying surface, etc.. After the surface geometry is created, the imported surface is segmented into sub-surfaces based on the dihedral angle between neighbouring facets. This results in creation of sub-surfaces which together form the imported surface. An example of a surface triangulation with some of its sub-surfaces is shown in Figure 1. The input geometry is segmented into ten sub-surfaces (four shown in the figure). The boundaries of these sub-surfaces are highlighted. These sub-surfaces are meshed separately by the advancing layer surface mesh algorithm.

The imported geometry in CGM is stored in memory while we create the surface mesh. The segmented surface gives us the freedom of meshing sub-surfaces independently. The ability of meshing sub-surfaces independently allows us to decide the anisotropy direction independently for the surfaces and can also be helpful if the surface mesh generation is required to run in parallel.

B. Boundary Initialization

After importing the surface triangulation, we have a valid underlying surface representation with us. Also, segmented sub-surfaces and their boundaries provide us with the boundary we need to march off of. The mesh generation routine starts by initializing the boundaries of each sub-surface of the surface by marking each point on the boundary as a candidate marching point which form the starting layer in the mesh. The points at the boundary of the mesh serve as the parent points for the first layer inserted into the mesh. The data structure for each boundary front vertex stores

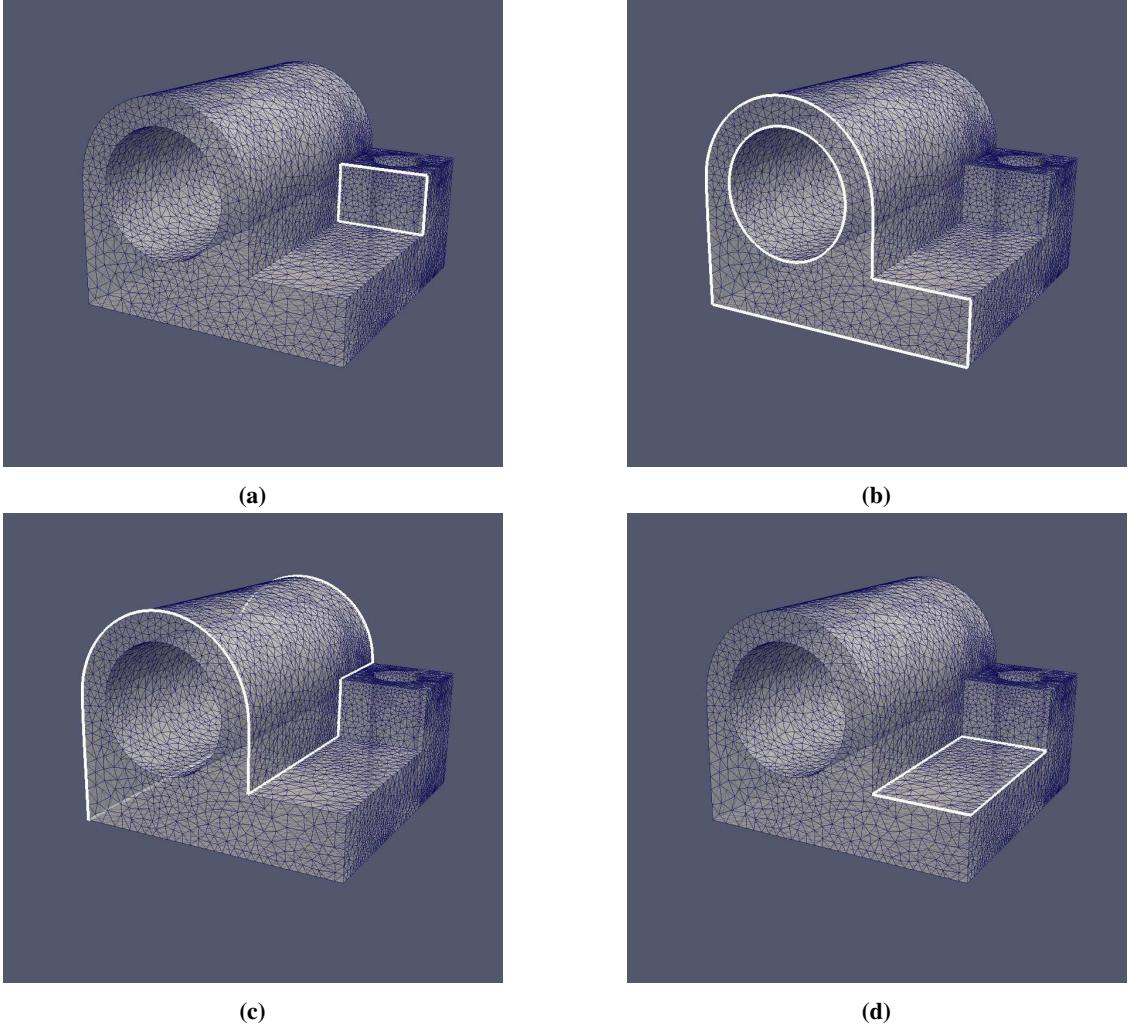


Fig. 1 An example input triangulation with some of the segmented surfaces.

its adjacent edges so as to identify the marching directions which is explained in detail in the next subsection. Each edge stores the direction into the interior of the sub-surface it bounds with respect to the surface normal. As the sub-segments sharing a common boundary are meshed independently, it is easier to identify the normal to an edge along the sub-surface for advancing the front in the mesh generation algorithm.

C. Advancing Layer Routine- Point placement

For each of the sub-surfaces of the geometry, the advancing layer routine iteratively picks a point from its boundary and extrudes it in a given direction. After evaluating the extruded point, we project the point onto the underlying surface. This process is set up to be of two steps for simplicity, accuracy and computational efficiency as will be explained later.

In the first step, we extrude the parent point to get the extruded point. We would interchangeably call the extruded points as the kid points as they represent the successors of their parent points from the previous layer. The direction of this extrusion is set up to be the average of the normals of the parent vertex's adjacent edges in the tangential plane of the sub-surface. In other words, the normals of the adjacent two edges of a vertex on the underlying sub-surface are averaged to get the extrusion direction. Consider Figure 2; we need to find the extrusion direction of a point P on the boundary. In the underlying triangulation T , PB and PA are the edges adjacent to P on the boundary of the surface S . We first find the points B_m and A_m , which are the mid-points of quartic-bezier curves PB and PA which are constructed by CGM as a part of constructing the quartic-bezier triangular surface patches from the underlying triangulation T . Next, we find the normal directions on the surface at point B_m and A_m which are labeled as N_b and N_a in the figure.

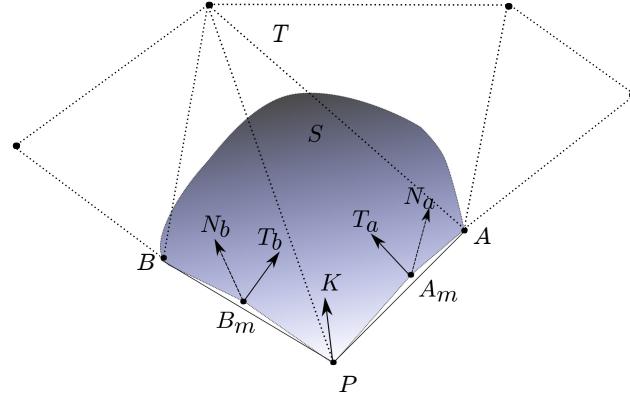


Fig. 2 Calculation of the extrude direction for a point P in the advancing front

We cross product the vector PB with N_b to get the direction T_b . The vector T_b is normal to the edge PB as well as tangential to the surface S . Similarly, we find the vector T_a . The direction of extrusion PK is chosen to be the average of the direction of T_a and T_b .

The initial extrusion length is an input parameter provided by the user. This length would be taken as the extrusion length when the boundary points are extruded for the first time to the interior of the surface. This length can vary with boundary vertices as the points are extruded independently. Hence, this extrude length can either be supplied by the user for all the points of the boundary separately or as a single value for all the boundary points. To get the best quality of quad elements, we are also experimenting with a varying extrude length along the boundary with respect to the interior angle between the direction vectors T_a and T_b . If the vertex is a concave corner vertex, the extrude length is increased so as to create good quality quad elements in the next layer and vice-versa for convex corner vertices. Also, if the vertex is a convex corner, additional marching directions are added so as to sufficiently refine the interior domain of the surface. The number of additional marching directions added depends on the convexity of the vertex. If we have n marching directions for a point, the angle between the adjacent edges would be divided into $n + 1$ angles to get the marching directions for the n points. We hope to include these modifications in the final paper submission.

After we have extruded the point, we project it on to the underlying geometry. This operation ensures that all the points we insert in the mesh are on the underlying geometry and avoid compounding the errors in subsequent layer generation process. Points are inserted in the mesh and the mesh elements are subdivided to include the new point. The candidate point for insertion can subdivide an existing triangle to replace the previous triangle with three new ones, or can subdivide an edge to replace existing two triangles with four new ones. To find the best triangle or edge for subdivision, we first make a guess for the triangle to insert the point. Any triangle in the surface interior adjacent to the point being extruded is chosen. Starting from this triangle, we iteratively jump to the best edge or triangle by comparing the barycentric coordinates of the new point with respect to the triangle in consideration. This technique suffers from two disadvantages. First, we need to compare double precision values of barycentric coordinates for making a decision on which triangle to choose for insertion. If the values are too close, the point might be inserted in the wrong triangle and would eventually lead to deviation of the mesh from the underlying surface. Second, the process of iteratively finding the right triangle for insertion might end up being in an infinite loop. Both of these problems are substantially reduced with a good isotropic initial triangulation. However, we add several validation tests to avoid these problems even for a coarse initial triangulation. These include orientation checks of the triangles formed with respect to the surface, thresholding the maximum deviation of the newly formed triangle from the surface and thresholding the dihedral angles between two triangles on the surface. We use an epsilon value of 10^{-5} while comparing the values of barycentric coordinates to zero. Also, we insert the point on a face rather than in a triangle when the ratio of the second-smallest barycentric coordinate to the smallest one is more than a set threshold(10^2). This helps us avoid very skinny triangles with large obtuse angles and also helps in avoiding several unnecessary face swapping in the mesh.

After advancing one layer to the surface interior, we increase the extrude length at each point by a factor. This factor, called the growth ratio would specify the anisotropic layer-on-layer extrusion length growth as we march on the surface. A value of growth ratio between 1.1 and 1.4 gives us satisfactory anisotropy at the boundaries of the mesh.

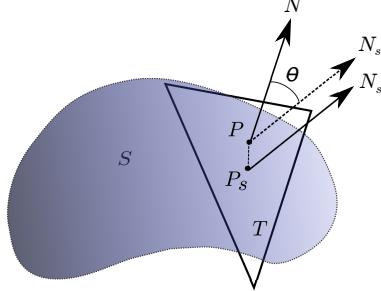


Fig. 3 Triangle T deviates from the underlying surface S . P is the centroid of T and P_s is the projection of P on S . The deviation is calculated as the interior angle between the normal to the triangle PN and the normal to the surface at P_s , which is P_sN_s . The angle θ represents the deviation here. The deviation is exaggerated for illustration purposes.

D. Local Reconnection for quality

A mesh observer object is initialized for the mesh data structure. The observer monitors all the changes happening to the mesh. These include creation and deletion of vertices, edges and cells in the mesh. An edge swapping manager and quality criterion for deciding a swap are also initialized for the mesh. The primary quality criterion we chose for an edge swap is maximization of the minimum angle in the pair of triangles adjacent to the edge before and after the swap. However, only this criterion is not enough for improving surface mesh quality. This is because a better quality pair of triangles may result by edge swapping according to this criterion, however, they might deviate quite a bit from the underlying surface. Hence, we threshold the maximum deviation of the resultant triangles from the surface. This deviation is calculated as the angle between the normal to the triangle and the normal at the centroid of the triangle projected onto the surface. See Figure 3 for reference. Deviation of a triangle T with respect to a surface S is shown. P is the centroid of the triangle and P_s is the projection of the centroid onto the surface. PN is the normal to the triangle and P_sN_s is the normal to the surface at P_s . The interior angle θ is considered as the deviation of the triangle from the surface. A upper bound is set for θ to limit the deviation of resultant triangles from the surface. The bound is set to 30° for the surfaces meshes we are dealing with. This bound is sensitive to the refinement in the initial triangulation. Basically, a higher value of θ would be apt for a coarse initial triangulation while the value can be set low for a fine initial triangulation.

We queue up the edges of the triangles affected by the insertion of the new point and do edge swapping for these edges to improve on the existing mesh. Figure 4 shows the process of point insertion and local reconnection in two steps. First, a point is inserted into the mesh at the desired location. This results in subdivision of a triangle as can be seen in figure 4b. Subsequently, edge swaps occur to improve the mesh quality. Two iterations of edge swap can be seen in figure 4c and 4d. We do not swap the edges which comprise the front or the edge between parent and child points. Doing so would invalidate the advancing front. Hence, only interior faces of the surface or the diagonals of the quad elements generated are swapped by the face swapping algorithm. After edge swapping, the algorithm moves on to the next point in the current layer to repeat the same process.

E. Front Recovery

In an advancing layer mesh generation algorithm, the accurate representation of the boundary by anisotropic layers generated from it is very important. Also, generation of quad-elements from triangles with the right aspect ratios is far easier if we retain the boundary representation as we advance multiple layers. Hence, after all points of the parent layer are extruded, we implement a front recovery routine which recovers all edges between kid vertices.

Front recovery is done by selecting a pair of points from the parent layer and then forcing an edge between the points extruded from these parent points; that is by forcing an edge between the kid vertices in the extruded layer. An edge is forced between the kid vertices by iteratively swapping any edge that lies topologically between the two kid points. This kind of forced swapping in the mesh might seem to decrease mesh quality. However, with the right boundary discretization, it would help in giving us the high-aspect ratio elements in the mesh that we desire. After all the kid vertices are connected by this method, validation checks are run to ensure that the extruded layer is well defined and there are no discrepancies in this layer as it will serve as the parent layer for the next one. This incorporates checking the connectivity of each vertex with its adjacent edges and the connectivity of each edge with its end points. Figure 5

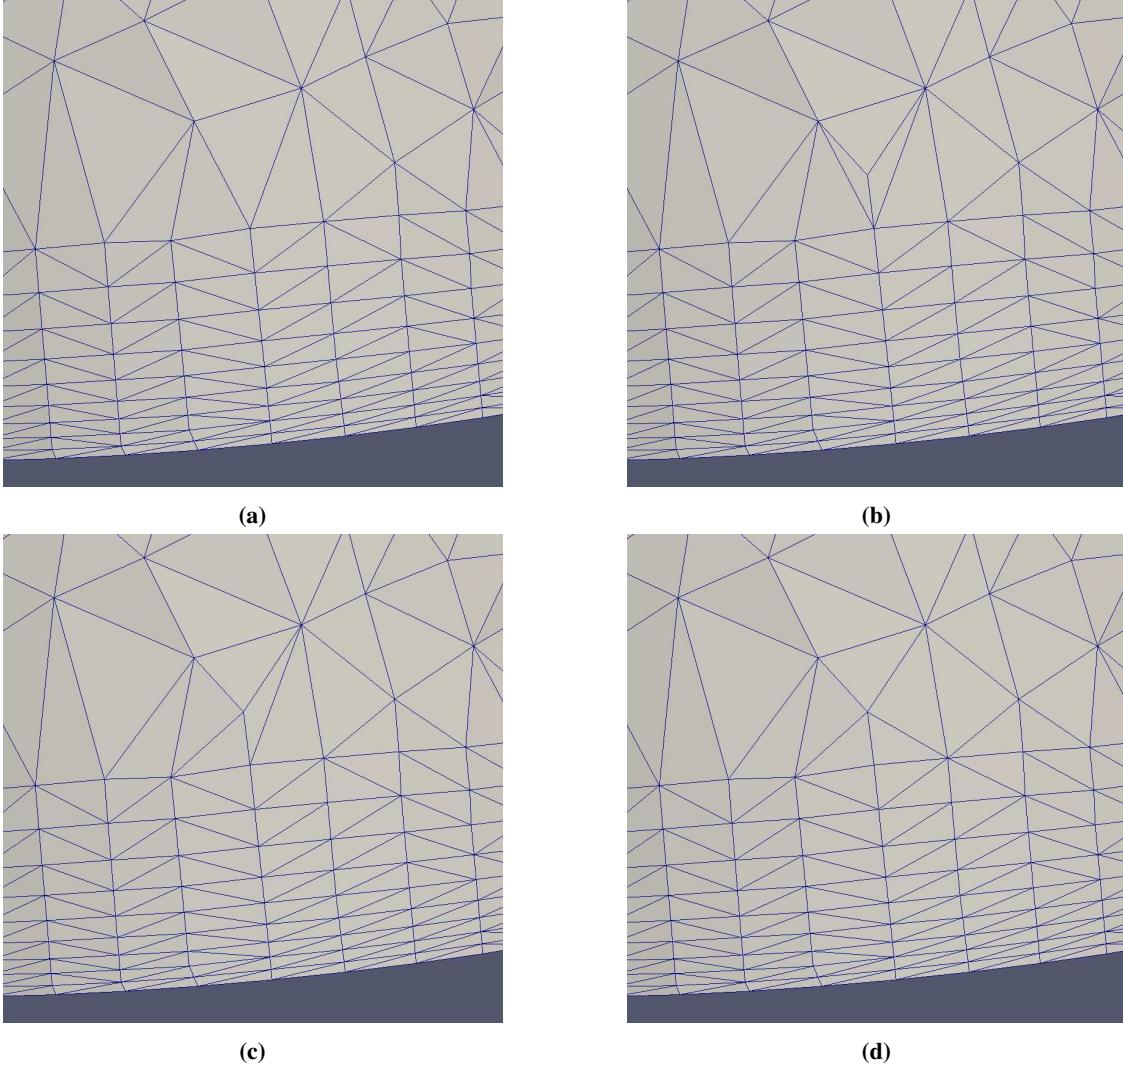


Fig. 4 Point Insertion and local reconnection for quality shown in four steps. (a) is the initial state of the mesh. A point is inserted in the mesh after extrusion from the parent point, projection onto the surface and finding the right triangle to insert. The triangle is subdivided into three new triangles as shown in (b). To improve the mesh quality after point insertion, swapping is done based on maximization of the minimum angle in a triangle. Two swaps occur as shown in figure (c) and (d) to improve mesh quality.

shows how exactly we recover an edge in the front by iteratively swapping the edges which obstruct the edge creation between two kid vertices, P_1 and P_2 in the figure. The edges which are topologically obstructing kid vertices P_1 and P_2 are swapped iteratively. A test is setup for each swap to avoid the flipping of triangles on the surface. After all the obstructing edges are swapped, a front edge between the kid vertices P_1 and P_2 would be recovered as shown in green in Figure 5e.

F. Edge Collapse and Interior Improvement

As the advancing front moves into surface interior, we decimate interior vertices of the initial triangulation as well as boundary vertices which are too close to each other through edge-collapse.

Once we have recovered the advancing front by iterative edge swaps to connect the kid points in the mesh, we check for vertices on the front that are too close to each other relative to the extrusion length. For instance, vertices which are near a concave corner could encroach each other if they are not decimated. Another example would be vertices on

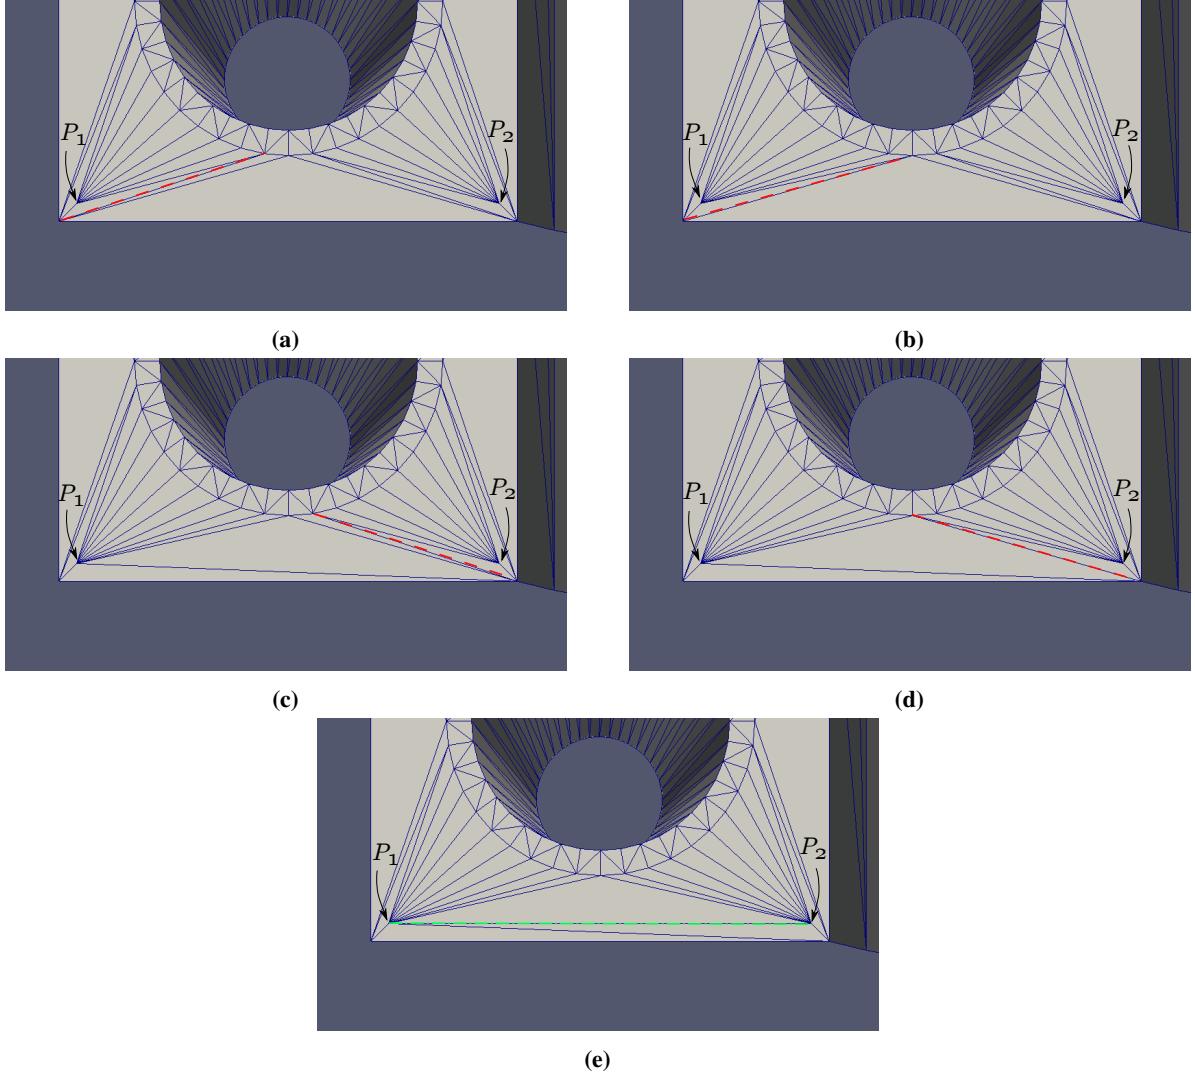


Fig. 5 Front Recovery: Point P_1 and point P_2 here represent the two kid points generated from the boundary of the surface. We try to force a connection between the two points by iteratively swapping edges which topologically obstruct their connection. Red dashed lines represents the edge chosen to be swapped next. In (e), the green edge is the edge recovered and would serve as a part of the next front in the mesh. Note that this example is for front recovery illustration purposes and the initial boundary discretization is too coarse to get a good-quality advancing layer mesh.

the advancing front which are far away from surface boundaries because the extrusion length has grown by this point. Decimating vertices from the front which are at a substantial distance from surface boundaries helps prevent the cell aspect ratio, ie front edge length over extrusion length from dropping below one.

To check for vertices to decimate, we iterate through the vertices in the front and identify the ones which are too close to their neighbours. Vertex decimation is done through the edge-collapse routine as described in [27]. The threshold edge length between two points on the front is set to be $2\tan(\pi/8)$ times the average extrusion length at those points. This value is set so as to minimize the normalized maximum deviation of angle from 90 degree for quad elements. Hence, the threshold ensures that the anisotropic properties are retained for several layers into the surface. All short edges on the front are collapsed using the edge-collapse algorithm. An example of an edge-collapse on the front is shown in Figure 6. Here, two points P_1 and P_2 are collapsed into a single point which forms a part of the next front on the surface.

As the advancing front marches into the surface, vertices in the interior of the surface immediately next to the front

are decimated to make way for the advancing layers. Before extruding a point on the advancing front, we check if any point in the surface interior with which it shares an edge agrees with the following condition. If it does, we decimate the interior vertex.

$$d < \max(l_1/\sqrt{2}, l_2/\sqrt{2}, c \times \text{extrudeLength}) \quad (1)$$

Here d is the distance between the point on the advancing front and the interior point, l_1 and l_2 are the lengths of adjacent edges of the point on the front, c is a constant whose value is set as $\sqrt{2}$ and extrudeLength is the extrusion length at the vertex on the front. This condition ensures decimation of vertices in surface interior which are close to the advancing front and avoids any encroachment of surface interior vertices on the advancing layers. Topological and geometrical checks are done before an edge can be collapsed. These include a threshold for the ratio of area of generated triangles and a limit on the dihedral angle between the adjacent triangles created by the collapse. The area threshold is set to be 10^8 . Also, edge-collapse is successful only if the triangles resulting from it are within a limit of $\theta < 30^\circ$ from the surface (see Figure 3). The threshold for the maximum dihedral angle between any two adjacent triangles that result from the edge-collapse is set to 40° . The quality criterion used for edge collapse is again maximization of the minimum angle in the triangles thus produced. The best edge for collapse is chosen when decimating the interior vertices using this quality criterion. This is in contrast to the vertex decimation on the advancing front where the candidate edge for collapse is already identified. An illustrative example for surface interior vertex decimation can be seen in Figure 7.

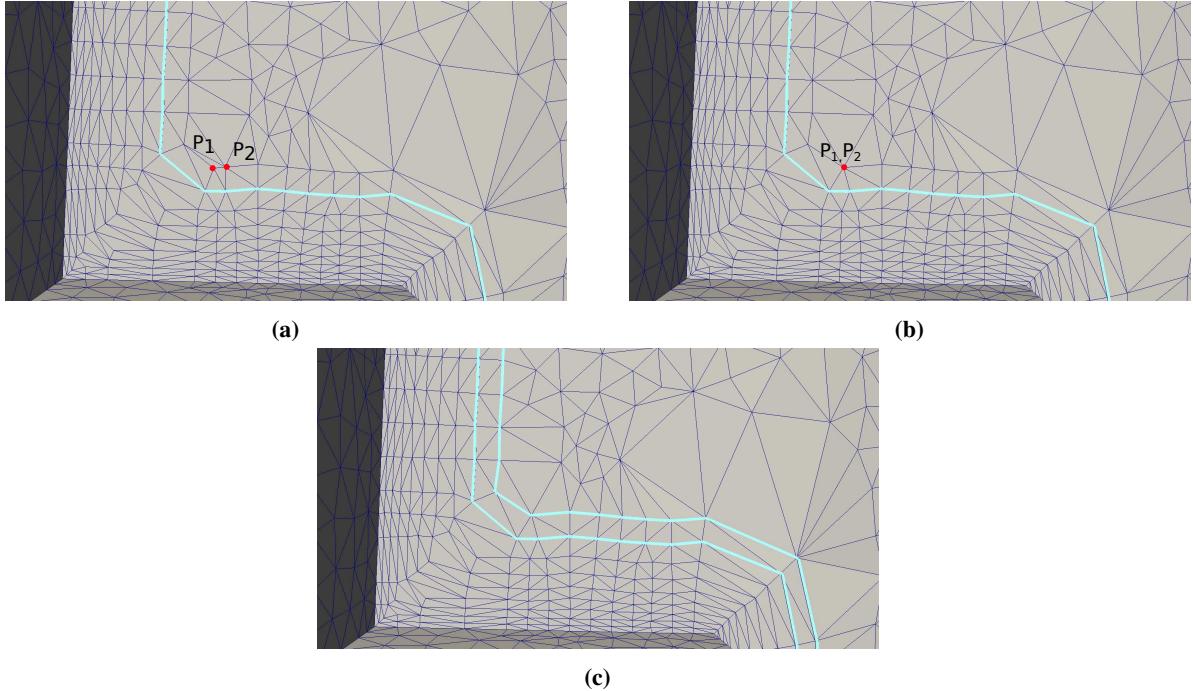


Fig. 6 Edge collapse on an advancing front to avoid encroachment of points. In (a), two points in the kid layer P_1 and P_2 are sufficiently close to each other. Their parent layer is highlighted. If both the points advance to the next layer, then the next front would fail to recover. Hence, the edge between them is chosen to collapse. (b) shows the result of the edge collapse. The new location of both the points is the average position of their initial location. (c) shows how the next front looks like.

After we have recovered the front and decimated encroaching vertices in the surface interior as well as on the advancing front through edge collapse, we queue up the immediate interior edges of the surface and swap them for maximizing mesh quality. This step is included so that we have a good interior triangulation at each step of the advancing layer routine. This step ensures that we do not have bad triangles causing problems as we continue to march.

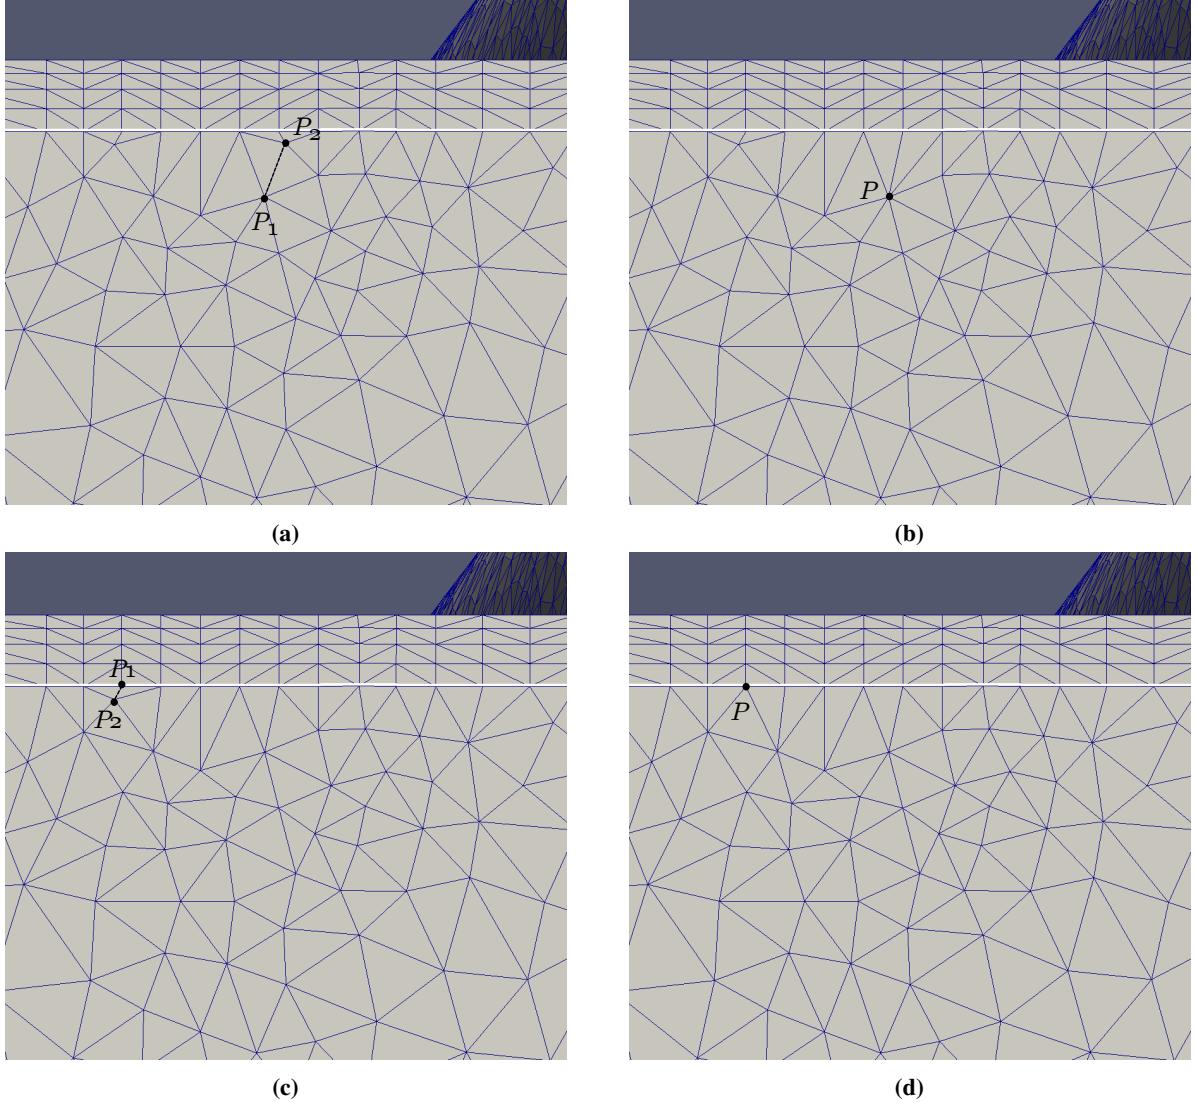


Fig. 7 Interior vertex decimation through edge collapse. The highlighted line shows the advancing front. In (a), vertex P_2 is about to encroach the front. Hence, the best vertex for collapse is chosen among its neighbours. The best vertex for edge collapse here is P_1 . Hence, P_2 is collapsed on to P_1 . The connectivity after the edge collapse is shown in (b) where vertex P represents the collapsed vertex. Similarly, in (c), vertex P_2 is about to encroach the advancing front and is collapsed onto vertex P_1 which is on the advancing front itself. The new connectivity is shown in (d) where all the possibly encroaching vertices for the advancing layer are decimated.

G. Front Collision and Redefinition

Till now we have described how points advance in layers and how are the layers recovered after each advancing layer step. However, some points on the advancing layer might encroach another edge on the advancing front. In other words, a point in the advancing front might come too close to other points(or edges) on the front which might result in encroachment if we advance all the points in the current layer to the next one. To tackle this issue, we preemptively identify such points and redefine the front so as to avoid encroachments. The redefinition process results in termination of the marching process for some or all the points on the advancing front.

To achieve the goal of front redefinition, we first iterate through all the points in the front to identify the encroached points. These points are the ones whose distance from any edge in the front, except its adjacent edges, is less than 2.4 times the extrusion length at that point. As our surface mesh is closed, we can utilize vertex connectivity to find out possible vertex-front edge encroachment. Hence, for a given point on the advancing front, only the edges on the

front that share a triangle with that point are checked for encroachment. This saves a lot of computational cost while identifying encroached vertices. This is assuming that the input triangulation is fine enough to resolve the complex features of the geometry so that we don't have to iterate through all the edges to find whether a given point is encroached.

After these encroached points and their corresponding encroached edges are identified, we reconnect the non-encroached points so as to form the new front. The encroached points are moved out of the advancing front and the front is hence redefined. An example of front redefinition can be seen in Figure 8. Here, the red line denotes the old front. If two of the three corner points are marched on to the next layer, the mesh elements would overlap each other and the front would be invalidated. Hence, two of them are selected to be moved out of the front. The front is redefined to contain the adjacent edges of only the remaining vertices in the front. The green line shows the newly defined front which would be used for the generation of the next layer on the surface.

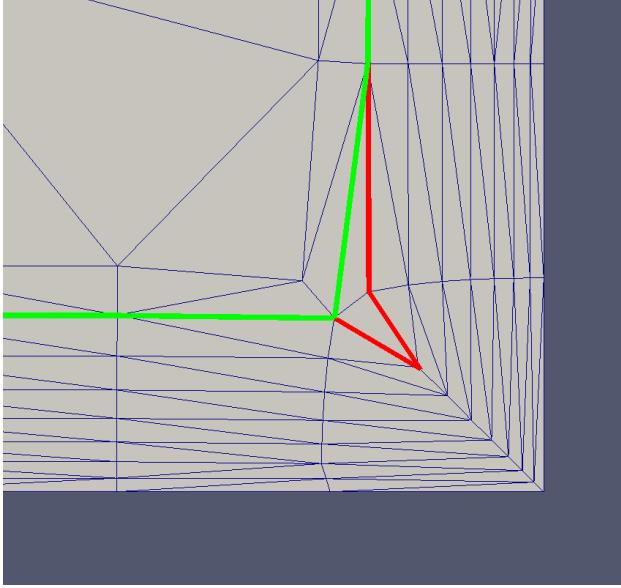


Fig. 8 Zoomed in view of an advancing layer mesh. After advancing a few layers, the mesh reaches the front denoted by red line. If the mesh keeps advancing like this, the front would collapse on itself resulting in undefined behaviour. The front redefinition routine stops two points from marching further and redefines the front as shown in green. The next marched layer is also shown.

H. Combining Triangular Elements to Quads

Generating advancing layers over the surface proceeds by creating triangular elements, which could be easily converted into quadrilateral elements.

III. Overall Mesh Generation Algorithm

In the previous section, we have detailed the steps in the advancing layer mesh generation routine. Here, we summarize the algorithm in the form of a pseudo code in Algorithm 1.

IV. Assumptions

We make a few assumptions on the input triangulation for our mesh generation algorithm to successfully create a valid mesh. The first one is that the triangulation of the surface is fine enough to resolve the geometric complexities in the object. If the input triangulation is too coarse, the encroaching interior vertex deletion subroutine will create mesh elements that deviate significantly from the surface or even fail. However, this issue can be tackled by refining the triangulation using any isotropic refinement algorithm as long as it retains the features of the input object. Coarse triangulations are acceptable for objects without any geometric complexities.

Another assumption we have made for the input triangulation is that it would have well defined boundaries where the

Algorithm 1 Overall Mesh Generation algorithm

```
1: procedure SURFMESH::CREATEMESH(triangulation  $T$ , extrusion length, growth ratio)
2:   Create surface geometry  $S \leftarrow T$ 
3:   Initialize Advancing Layer  $F$  from Surface Boundary
4:   while  $F \neq \emptyset$  do                                 $\triangleright$  Advance until Front  $F$  isn't empty
5:     ADVANCELAYER( $F$ )
6:     Validate the new Layer
7:   Export SURFMESH
8: procedure SURFMESH::ADVANCELAYER( $F$ )
9:   Delete Encroaching Interior Vertices
10:  Redefine front  $F$  Until No Vertex Encroachment
11:  Extrude All Points of  $F$  and Insert Into Mesh
12:  Recover front  $F'$                                  $\triangleright$  Through forced edge creation between kid vertices
13:  Improve Mesh Quality by Edge Swapping            $\triangleright$  In the immediate interior of mesh
14:  Collapse Short Edges in the Front
15:  Update Extrusion Length for vertices
```

advancing layers can march out from. For future, we intend to incorporate a subroutine that identifies these boundaries even on objects which don't have well defined features, for eg. leading edge of an airfoil. Accepting boundary elements as a user input is also an option for such objects.

V. Examples

We use the surface mesh generation algorithm described here to mesh some geometries. We illustrate how the advancing layer algorithm works by creating several advancing fronts from the boundary of input triangulation. Figure 9a shows a cylindrical object with a coarse boundary and surface discretization. The object contains three sub-surfaces. These include the planar end caps of the cylinder and the curved cylindrical surface. The three sub-surfaces of the input triangulation are segmented and meshed independently using the advancing front surface generation routine. Resultant mesh is shown in Figure 9b. It can be seen that the boundary curves, that is the circular bounding curve of the sub-surfaces is preserved very well for several layers in the surface interior. Hence, the advancing layer surface mesh generation algorithm helps to preserve the boundary curve even for coarse initial surface discretization. Figure 9c and 9d show the initial triangulation and the final mesh for the end cap of the cylinder.

Another example is shown in Figure 10. Here, the same cylindrical geometry is taken as the input but with a finer boundary and surface discretization as shown in Figure 10a. The surface mesh generated for this geometry is shown in Figure 10b. Here again, the advancing layers grow anisotropically from the boundary into the interior of the mesh and generate 16 layers on all the three sub-surfaces. The parameters used for both the meshes described for the cylindrical geometry are shown in Table 1. The value of aspect ratio at the surface boundary is set to $\approx 5 - 6$ to generate the anisotropic meshes. This value can be modified by changing the initial extrusion length and the growth ratio, which are the inputs to the algorithm. The number of layers generated on all three segmented surfaces of the coarse cylindrical geometry are 9 and for the fine cylindrical geometry are 16.

We demonstrate two more examples of advancing layer surface mesh generation. The first one is shown in Figure 11. Here, one of the sub-surfaces of the geometry shown previously in Figure 1 is meshed using the advancing layer surface mesh generation routine. The resultant mesh is shown in Figure 11b and 11d. The sub-surface boundary is highlighted in the figure. We can see that the regions close to surface boundaries have the required stretched elements having high aspect ratio. The aspect ratio for this mesh at the boundary varies with the boundary vertices but is set to be approximately 4-8. This gives us directional anisotropy as we move to surface interior from the boundaries of the mesh. Another example is shown in Figure 12. Here, a twisted object with a coarse initial triangulation is input to the surface mesh generator. The object is segmented into three sub-surfaces which are meshed independently. We advance seven layers from the sub-surface boundaries to get a mesh which can be seen in Figure 12b. A close up view on the mesh elements is shown in Figure 12d. Here, the aspect ratio of the elements next to sub-surface boundaries is set to be approximately 12-20. These examples demonstrate the capability of the surface mesh generator to successfully advance multiple layers of anisotropic elements on a surface while preserving the boundary features.

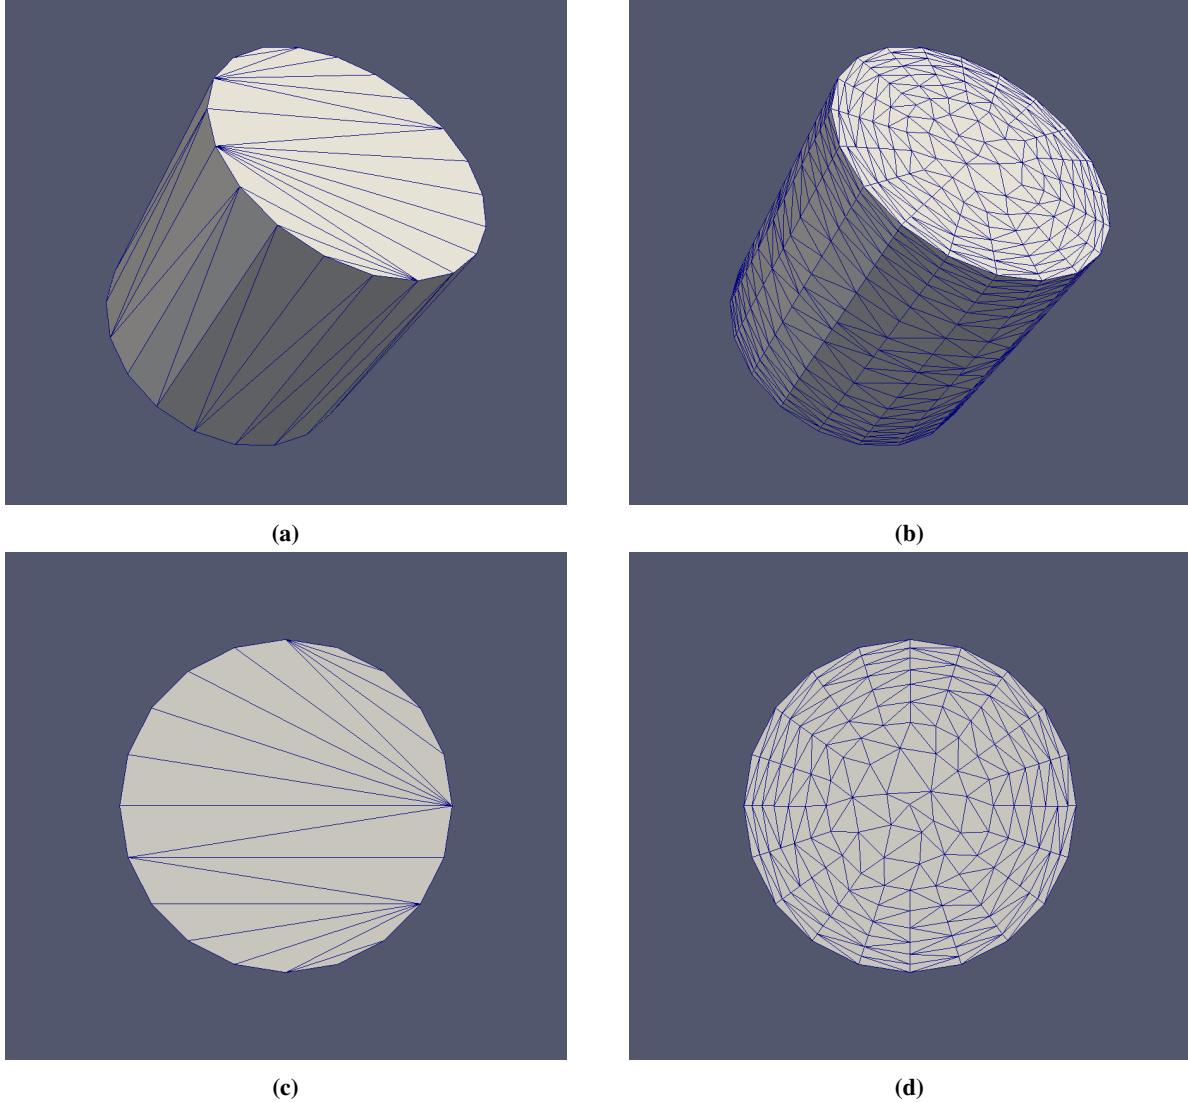
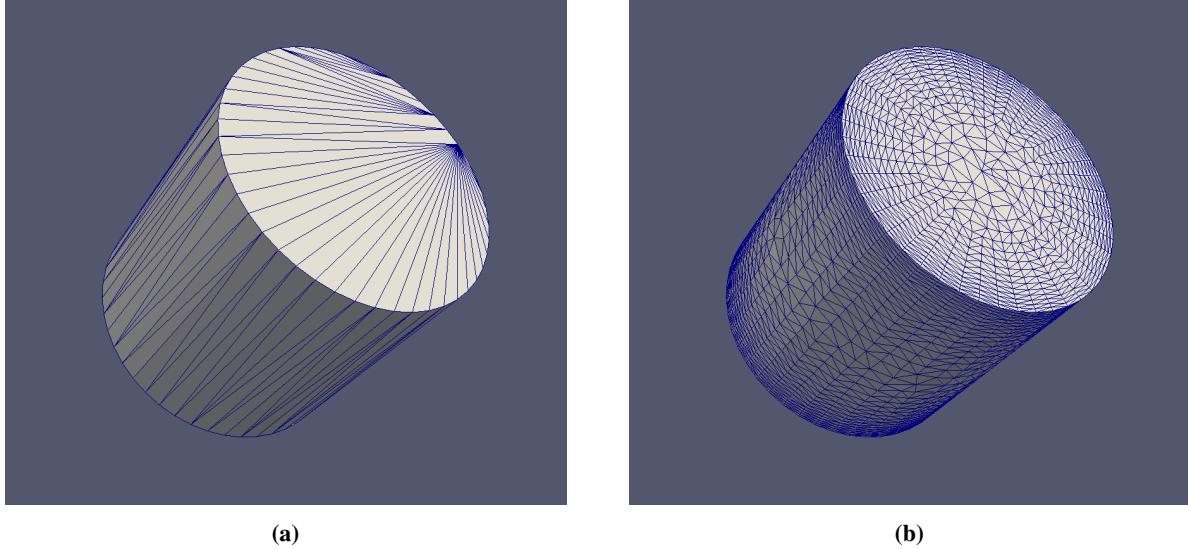


Fig. 9 Advancing layer mesh for a cylindrical object. (a) and (c) show the isometric view and the end cap of the initial surface discretization of the object. (b) and (d) show the surface mesh produced by the advancing layer routine.

VI. Work Plan for Complete Paper

In the final paper, we would present examples of completed surface meshes for more complex geometries. The following objectives would be met to achieve that.

- Merge pairs of triangles to form quads. This would be relatively easy as a front edge with a proper descent defines a quad.
- Add more than one extrusion points to convex corners of the geometry to make the initial aspect ratio of the advancing front more uniform.
- Add robustness to the front collision and redefinition routine so as to handle geometries with complex features and sharp concave corners.
- Identify or take user input for the direction of anisotropy at each extrusion point on the advancing layer.
- Add smoothing based on spring force model. The smoothing would help to ensure that the aspect ratio of the mesh elements remain consistent along the advancing front. It would also ensure that the mesh elements generated deviate as little from the underlying surface as possible.
- Vary extrusion length at a point with the angle between its adjacent edges on the advancing front



(a)

(b)

Fig. 10 Advancing layer mesh produced for a cylindrical object with finer boundary discretization. The boundary curve is preserved as the advancing layers march on the surface of the object.

Table 1 Mesh parameters for Cylindrical Surface Mesh

Case	Edge length	Initial Extrusion Length	Growth Ratio	Aspect Ratio	No. of layers
Coarse Boundary	3.06	0.50	1.2	6.14	9
Fine Boundary	1.34	0.24	1.1	5.57	16

References

- [1] Sbardella, L., Sayma, A., and Imregun, M., “Semi-structured meshes for axial turbomachinery blades,” *International Journal for Numerical Methods in Fluids*, Vol. 32, No. 5, 2000, pp. 569–584.
- [2] Dyedov, V., Einstein, D. R., Jiao, X., Kuprat, A. P., Carson, J. P., and Del Pin, F., “Variational generation of prismatic boundary-layer meshes for biomedical computing,” *International Journal for Numerical Methods in Engineering*, Vol. 79, No. 8, 2009, pp. 907–945.
- [3] Mavriplis, D. J., “Adaptive mesh generation for viscous flows using triangulation,” *Journal of computational Physics*, Vol. 90, No. 2, 1990, pp. 271–291.
- [4] Castro-Díaz, M., Hecht, F., Mohammadi, B., and Pironneau, O., “Anisotropic unstructured mesh adaption for flow simulations,” *International Journal for Numerical Methods in Fluids*, Vol. 25, No. 4, 1997, pp. 475–491.
- [5] Nakahashi, K., “FDM-FEM zonal approach for viscous flow computations over multiple-bodies,” *25th AIAA Aerospace Sciences Meeting*, 1987, p. 604.
- [6] Kallinderis, Y., Khawaja, A., and McMorris, H., “Hybrid prismatic/tetrahedral grid generation for viscous flows around complex geometries,” *AIAA journal*, Vol. 34, No. 2, 1996, pp. 291–298.
- [7] Pirzadeh, S., “Unstructured viscous grid generation by the advancing-layers method,” *AIAA journal*, Vol. 32, No. 8, 1994, pp. 1735–1737.
- [8] Löhner, R., “Matching semi-structured and unstructured grids for Navier-Stokes calculations,” *11th Computational Fluid Dynamics Conference*, 1993, p. 3348.
- [9] Connell, S. D., and Braaten, M. E., “Semistructured mesh generation for three-dimensional Navier-Stokes calculations,” *AIAA journal*, Vol. 33, No. 6, 1995, pp. 1017–1024.
- [10] Garimella, R. V., and Shephard, M. S., “Boundary layer mesh generation for viscous flow simulations,” *International Journal for Numerical Methods in Engineering*, Vol. 49, No. 1-2, 2000, pp. 193–218.

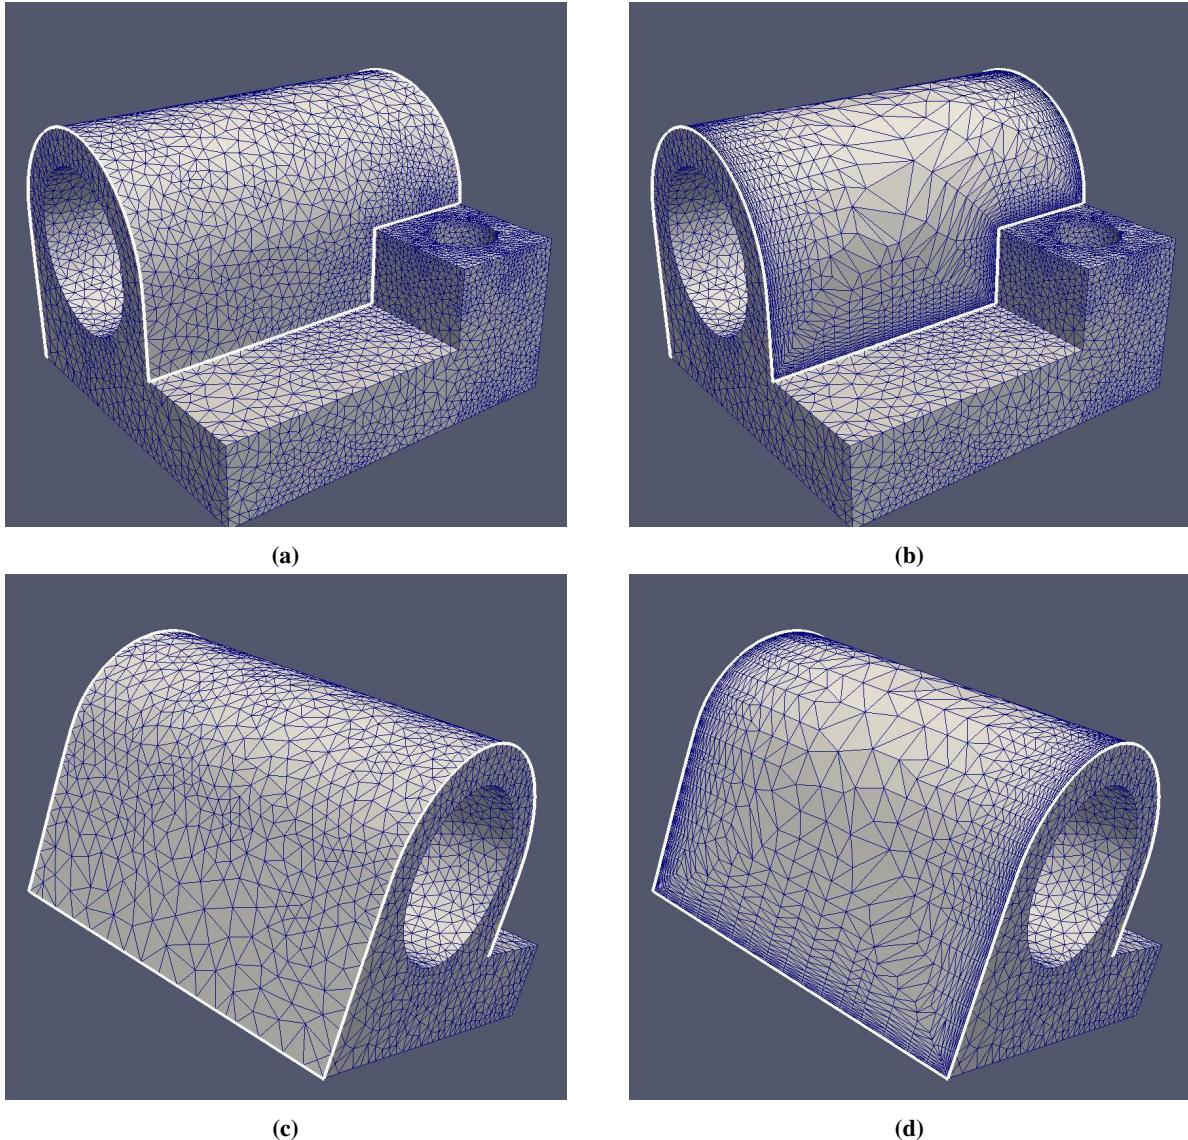


Fig. 11 One of the surfaces of the geometry remeshed with the advancing layer surface mesh generation algorithm. The boundary of the surface is highlighted.

- [11] Ito, Y., and Nakahashi, K., “Unstructured Mesh Generation for Viscous Flow Computations.” *IMR*, 2002, pp. 367–377.
- [12] Sahni, O., Jansen, K. E., Shephard, M. S., Taylor, C. A., and Beall, M. W., “Adaptive boundary layer meshing for viscous flow simulations,” *Engineering with Computers*, Vol. 24, No. 3, 2008, p. 267.
- [13] Ito, Y., Shih, A. M., Soni, B. K., and Nakahashi, K., “Multiple marching direction approach to generate high quality hybrid meshes,” *AIAA journal*, Vol. 45, No. 1, 2007, pp. 162–167.
- [14] Aubry, R., and Löhner, R., “Generation of viscous grids at ridges and corners,” *International journal for numerical methods in engineering*, Vol. 77, No. 9, 2009, pp. 1247–1289.
- [15] Mavriplis, D., “Unstructured grid techniques,” *Annual Review of Fluid Mechanics*, Vol. 29, No. 1, 1997, pp. 473–514.
- [16] Aftosmis, M., Gaitonde, D., and Tavares, T. S., “On the accuracy, stability, and monotonicity of various reconstruction algorithms for unstructured meshes,” 1994.

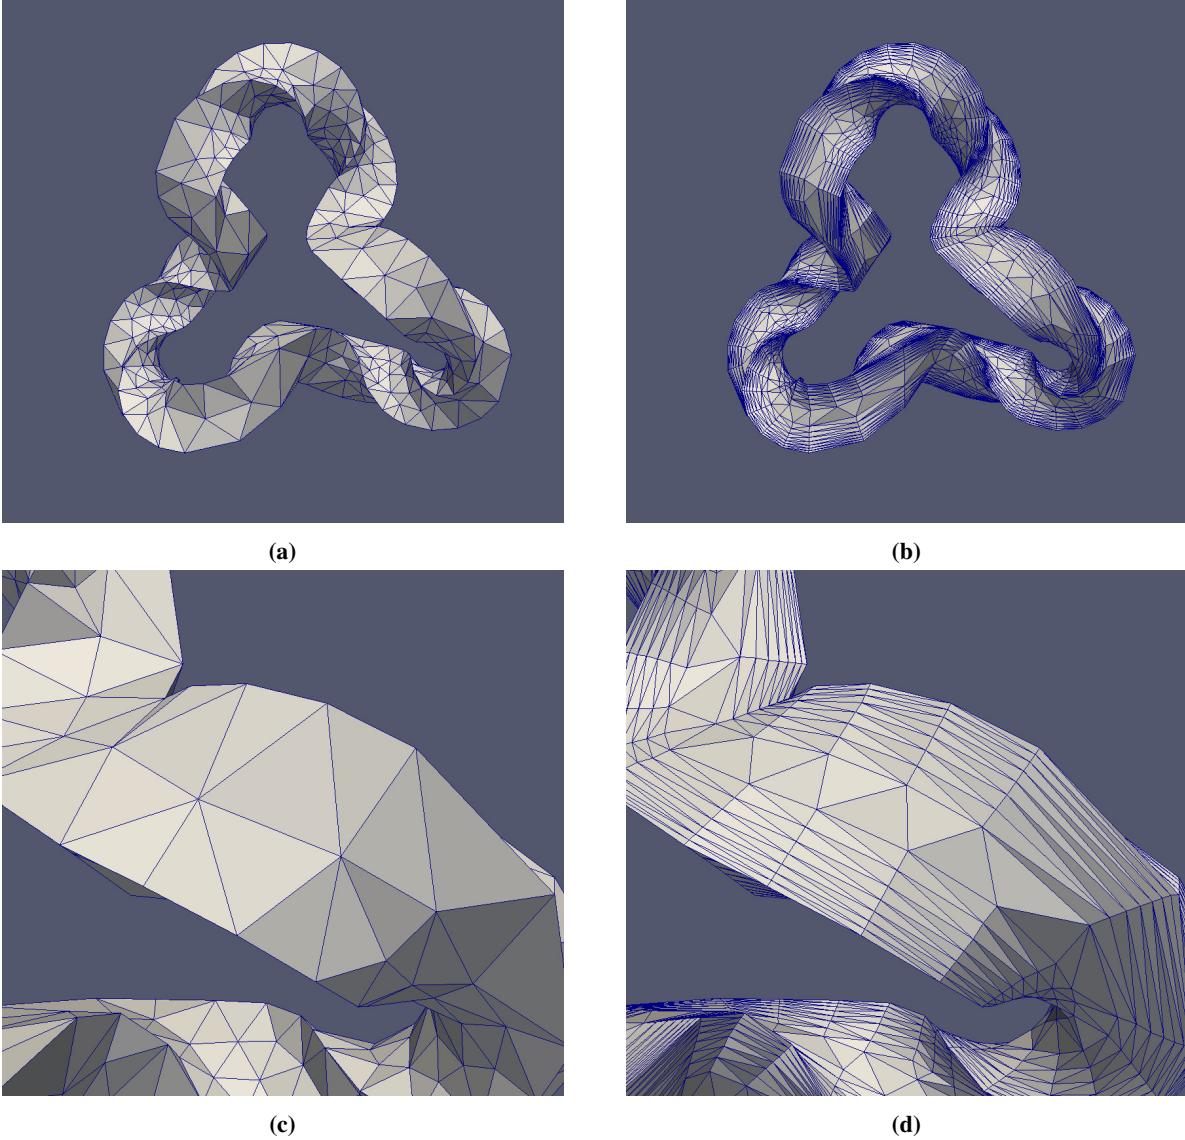


Fig. 12 A contorted initial geometry is meshed using the advancing layer algorithm for seven layers. The boundaries of the geometry are preserved even after advancing for several layers into the surface of the object.

- [17] Blacker, T. D., and Stephenson, M. B., “Paving: A new approach to automated quadrilateral mesh generation,” *International Journal for Numerical Methods in Engineering*, Vol. 32, No. 4, 1991, pp. 811–847.
- [18] Zhu, J., Zienkiewicz, O., Hinton, E., and Wu, J., “A new approach to the development of automatic quadrilateral mesh generation,” *International Journal for Numerical Methods in Engineering*, Vol. 32, No. 4, 1991, pp. 849–866.
- [19] George, P.-L., and Borouchaki, H., “Delaunay triangulation and meshing,” 1998.
- [20] Chen, H., and Bishop, J., “Delaunay triangulation for curved surfaces,” *Meshing Roundtable*, 1997, pp. 115–127.
- [21] Owen, S. J., “A survey of unstructured mesh generation technology.” *IMR*, 1998, pp. 239–267.
- [22] Cuillière, J.-C., “An adaptive method for the automatic triangulation of 3D parametric surfaces,” *Computer-aided design*, Vol. 30, No. 2, 1998, pp. 139–149.
- [23] Tristano, J. R., Owen, S. J., and Canann, S. A., “Advancing front surface mesh generation in parametric space using a riemannian surface definition.” *IMR*, 1998, pp. 429–445.

- [24] Lan, T., and Lo, S., “Finite element mesh generation over analytical curved surfaces,” *Computers & Structures*, Vol. 59, No. 2, 1996, pp. 301–309.
- [25] Lau, T., Lo, S., and Lee, C., “Generation of quadrilateral mesh over analytical curved surfaces,” *Finite Elements in Analysis and Design*, Vol. 27, No. 3, 1997, pp. 251–272.
- [26] Tautges, T. J., “The Common Geometry Module (CGM).” Tech. rep., Sandia National Laboratories, 2004.
- [27] Hoppe, H., DeRose, T., Duchamp, T., McDonald, J., and Stuetzle, W., “Mesh optimization,” Tech. rep., Washington Univ Seattle Dept OF Computer Science And Engineering, 1994.